



Game Theory

A red die with yellow pips is shown on a green game board. The game board has a blue number 9 and a yellow sheep icon. The background is a blurred landscape with green hills and a blue sky.

Motivation

- examine the role of AI methods in games
- some game provide challenges that can be formulated as abstract competitions with clearly defined states and rules
 - programs for some games can be derived from search methods
 - narrow view of games
- games can be used to demonstrate the power of computer-based techniques and methods
- more challenging games require the incorporation of specific knowledge and information
- expansion of the use of games
 - from entertainment to training and education



Objectives

- explore the combination of AI and games
- understand the use and application of search methods to game programs
 - apply refined search methods such as minimax to simple game configurations
 - use alpha-beta pruning to improve the efficiency of game programs
 - understand the influence of chance on the solvability of chance-based games
- evaluation of methods
 - suitability of game techniques for specific games
 - suitability of AI methods for games

Games and Computers

games offer concrete or abstract competitions

- “I’m better than you!”

some games are amenable to computer treatment

- mostly mental activities
- well-formulated rules and operators
- accessible state

others are not


- emphasis on physical activities
- rules and operators open to interpretation
 - need for referees, mitigation procedures
- state not (easily or fully) accessible

Defining Computer Games



video game / computer game

“a mental contest, played with a computer according to certain rules for amusement, recreation, or winning a stake” [Zyda, 2005]



serious game

a mental contest, played with a computer in accordance with specific rules, that uses entertainment to further training or education [modified from Zyda, 2005]

[Zyda, 2005] Michael Zyda, “From Visual Simulation to Virtual Reality to Games,” *IEEE Computer*, vol. 39, no. 9, pp. 25-32.

A decorative background on the left side of the slide. It features a close-up of two dice, one yellow and one blue, resting on a surface that appears to be a laptop keyboard. The background is blurred, showing warm colors like orange and yellow. A blue and grey geometric shape, resembling a stylized laptop or a corner, frames the right side of the text area.

Aspects of Computer Games

- story
 - defines the context and content of the game
- art
 - presentation of the game to the user
 - emphasis on visual display, sound
- software
 - implementation of the game on a computer
- purpose
 - entertainment
 - training
 - education



Games and AI

- traditionally, the emphasis has been on a narrow view of games
 - formal treatment, often as an expansion of search algorithms
- more recently, AI techniques have become more important in computer games
 - computer-controlled characters (agents)
 - more sophisticated story lines
 - more complex environments
 - better overall user experience
- Games (especially chess) have been a great focus for AI research. Why?
 - Uncertainty makes them like the real world
 - The rules can be well defined
 - The game position is easily represented
 - There are no moral or ethical problems
 - We can match computers against both other computers and humans
 - Playing games is seen as an "intelligent activity"



Cognitive Game Design

- story development
 - generation of interesting and appealing story lines
 - variations in story lines
 - analysis of large-scale game play
- character development
 - modeling and simulation of computer-controlled agents
 - possibly enhancement of user-controlled agents
- immersion
 - strong engagement of the player's mind
- emotion
 - integration of plausible and believable motion in characters
 - consideration of the user's emotion
- pedagogy
 - achievement of “higher” goals through entertainment

Game Analysis

- often deterministic
 - the outcome of actions is known
 - sometimes an element of chance is part of the game
 - e.g. dice
- two-player, turn-taking
 - one move for each player
- zero-sum utility function
 - what one player wins, the other must lose
- often perfect information
 - fully observable, everything is known to both players about the state of the environment (game)
 - not for all games
 - e.g. card games with “private” or “hidden” cards
 - Scrabble

Games as Adversarial Search

- many games can be formulated as search problems
- the zero-sum utility function leads to an adversarial situation
 - in order for one agent to win, the other necessarily has to lose
- factors complicating the search task
 - potentially huge search spaces
 - elements of chance
 - multi-person games, teams
 - time limits
 - imprecise rules

Difficulties with Games

- games can be very hard search problems
 - yet reasonably easy to formalize
 - finding the *optimal* solution may be impractical
 - a solution that beats the opponent is “good enough”
 - unforgiving
 - a solution that is “not good enough” not only leads to higher costs, but to a loss to the opponent
- example: chess
 - size of the search space
 - branching factor around 35
 - about 50 moves per player
 - about 35^{100} or 10^{154} nodes
 - about 10^{40} *distinct* nodes (size of the search graph)



Search Problem Formulation

- initial state
 - board, positions of pieces
 - whose turn is it
- successor function (operators)
 - list of (*move*, *state*)
 - defines the legal moves, and the resulting states
- terminal test
 - also called goal test
 - determines when the game is over
 - calculate the result
 - usually win, lose, draw, sometimes a score (see below)
- utility or payoff function
 - numeric value for the outcome of a game

A red die with yellow pips is shown on a green game board. The board has a blue number 9 and some yellow circular markers. The background is a blurred landscape with green hills and a blue sky.

Single-Person Game

- conventional search problem
 - identify a sequence of moves that leads to a winning state
 - examples: Solitaire, dragons and dungeons, Rubik's cube
 - little attention in AI
- some games can be quite challenging
 - some versions of solitaire
 - a heuristic for Rubik's cube was found by the Absolver program

Contingency Problem

- uncertainty due to the moves and motivations of the opponent
 - tries to make the game as difficult as possible for the player
 - attempts to maximize its own, and thus minimize the player's utility function value
- different from contingency due to neutral factors, such as
 - chance
 - outside influence



Two-Person Game

- How do we think when we play chess?
- We consider making a move....
 - *if I move my queen there, then my opponents' best move is to move their knight there, etc. etc.*
- We are making some assumptions:
 - we want to make our best possible move
 - our opponent is as skillful as we are
 - our opponent has the same information as us
 - our opponent will also do their best to win

Two-Person Game

- games with two opposing players
 - often called MIN and MAX
 - usually MAX moves first, then they take turns
 - in game terminology, a *move* comprises one step, or *ply* by each player
- MAX wants a strategy to find a winning state
 - no matter what MIN does
- MIN does the same
 - or at least tries to prevent MAX from winning
- full information
 - both players know the full state of the environment
- partial information
 - one player only knows part of the environment
 - some aspects may be hidden from the opponent, or from both players

Perfect Decision in Games

- Consider two players of a game, MAX and MIN
 - MAX moves first
- The game begins in an *initial state*
- There is a set of *operators*, i.e. legal moves
- *Terminal states*, where the game ends
- each terminal state has a *utility function*, i.e. a pay-off to each player
- The **Utility** is assumed to be in relation to MAX
 - we assume the Utility function(Evaluation function) is symmetrical, i.e. what is good for MAX is equally bad for MIN and vica-versa
- Examples of utility
 - chess might have a utility of 1 for a win, 0 for a draw and -1 for a loss
 - Zero sum games

Simple Games

- Assume that we can generate and search the entire decision tree
 - this is only possible for simple games, later we'll think about more complex games
- How should MAX move?
 - some lines end in wins for MAX, some in wins for MIN, but we don't know how MIN will move...

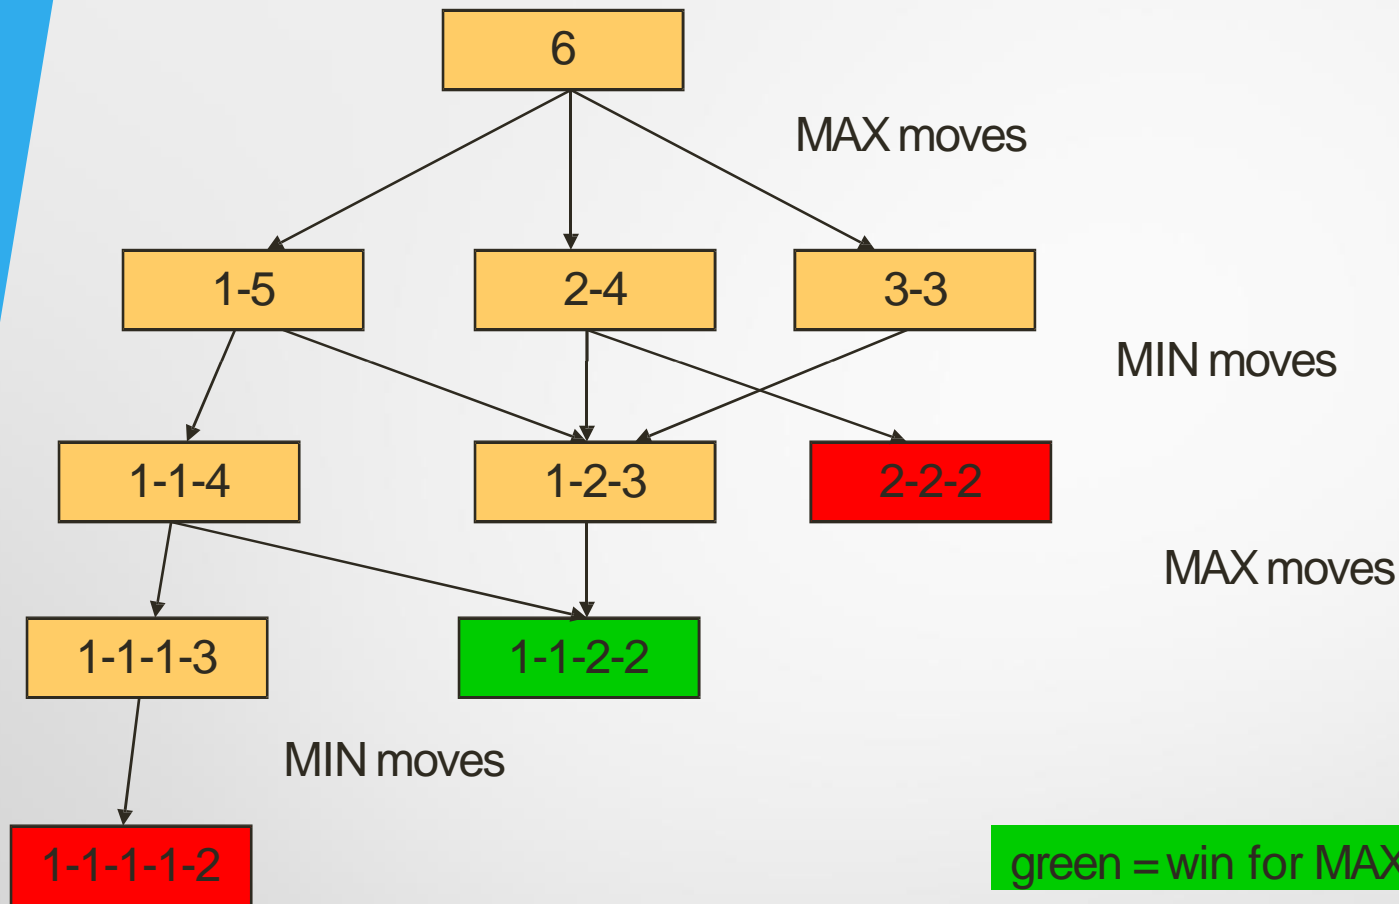
MINIMAX Searching

- What we do is first build the game tree
- Then we work in a bottom-up fashion, starting down at the terminal states
- What we assume is as follows:
 - When MAX is moving, MAX will chose the line with the highest utility, so pass up the MAXimum utility to the next highest level
 - When MIN is moving, MIN will chose the line with the lowest utility (for MAX), so pass up the MINimum utility to the next highest level
- When we reach the current position, MAX choses the line of highest utiltlity (as usual) and moves

Example of MinIMax

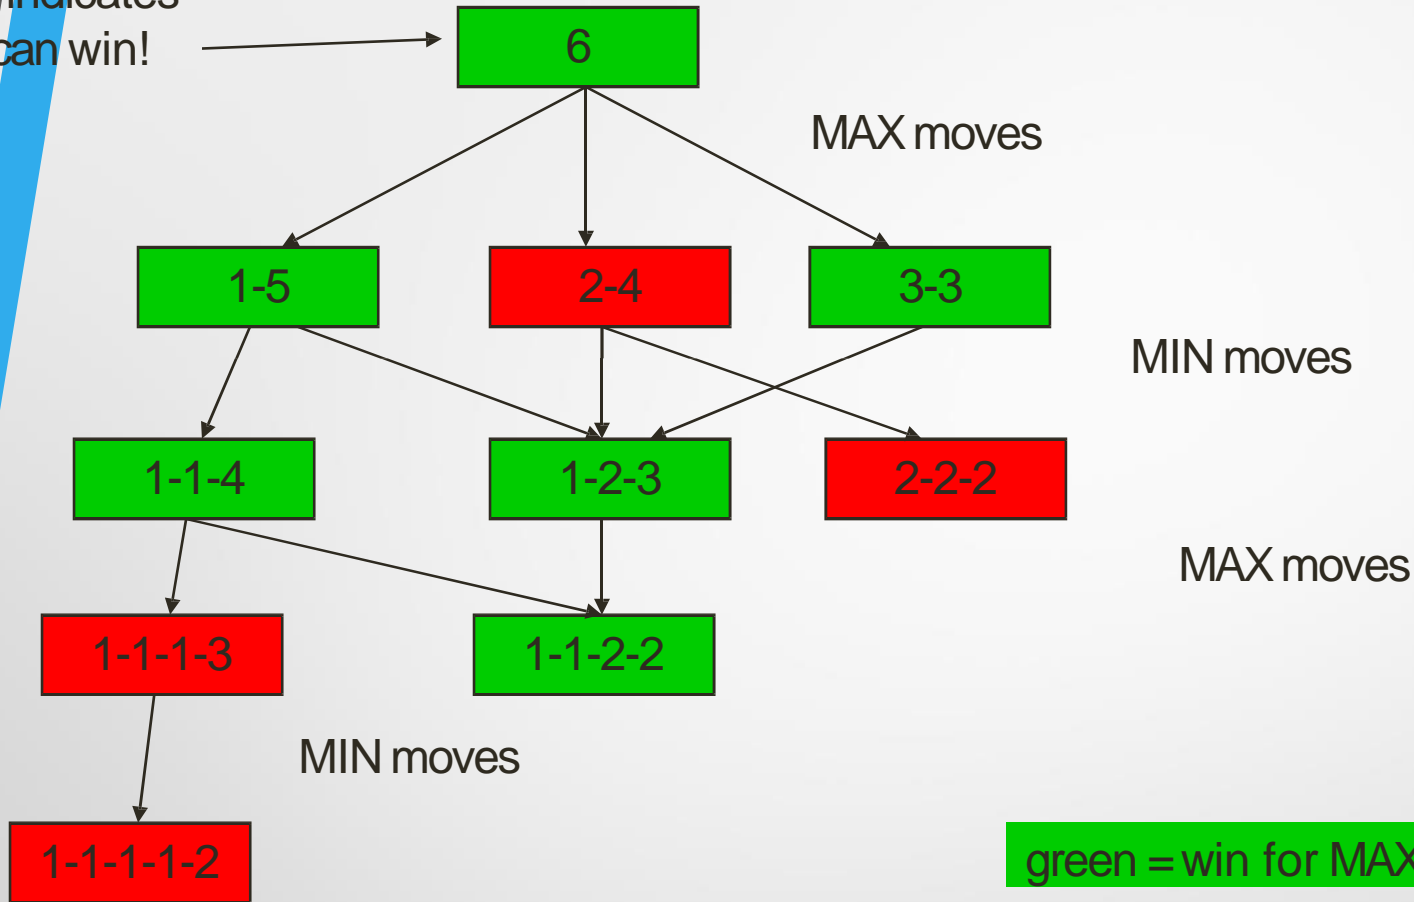
- A simple game is the game of *nim*
- The rules of *nim* are as follows:
 - we start with a number of sticks in a group
 - on each move, a player must divide a group into two smaller groups
 - groups of one or two sticks can't be divided
 - the last player who makes a legal move wins

NIM Search Tree



MINIMAX Search of NIM

green indicates
MAX can win!




Heuristic MINIMAX

- If the search tree is too large, we can't go down to the terminal states
- In this case we can search down only a certain number of levels (this is called the *PLY* of the search)
- At the maximum depth, we use a heuristic *evaluation function* to rate the utility of the different positions
 - then use the MINIMAX as usual to decide what to do

Practical Aspects of heuristic MINIMAX

- When we take each step, we will only need to calculate the evaluation function for the next level, provided we have stored the tree
 - this will save time but costs space
- Usually we search the tree depth-first
- We need to trade off sophistication of heuristics vs cost of searching

Heuristic MINIMAX

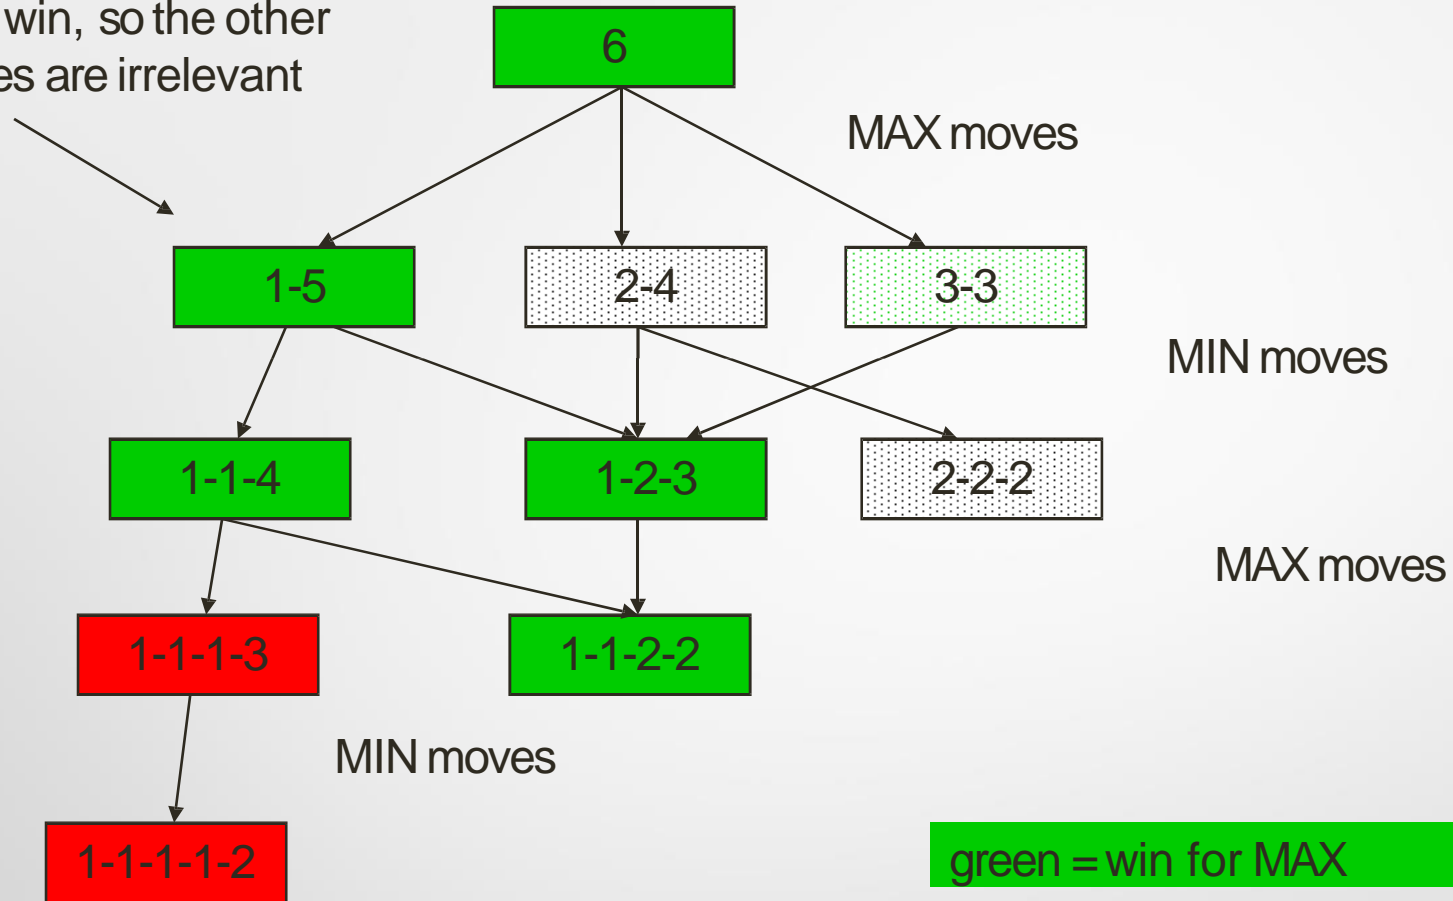


PRUNING

- When carrying out MINIMAX, we can save a lot of work without loss of performance by *pruning*
- Remember we are searching the tree depth first:
 - so in an instance where we have a simple discrete utility function, as soon as we have found a way to force a win, we don't need to look any longer
 - this applies at all levels when MAX is to move - as soon as we find a guaranteed winning line, we can prune other options
 - similarly, when MIN moves, we can prune if we find a guaranteed losing line (i.e. a win for MIN)


Pruning the NIM Search

We have found a certain win, so the other branches are irrelevant



Alpha-Beta Pruning


- We can generalise pruning for *continuous* utility functions through a procedure known as α - β
- During the depth first search point we remember:
 - the score for the best choice so far for MAX = α
 - the score for the best choice so far for MIN = β
- When evaluating moves for MAX, if we reach a point in the tree where MIN can choose a move with utility $\leq \alpha$, we can prune the move above
 - MAX will not choose this move but prefer the chain to where α was found (unless a better branch exists)
 - similarly for MIN



MAX


MIN

Alpha Beta Pruning



Alpha Beta Pruning


1. Start at C. Descend to full-ply depth and assign the heuristic to a state and all siblings (MIN 2, 3). Back up these values to their parent node (MAX 3).
2. Offer this value to the grandparent (A), as its beta value. So, **A has beta=3**. A will be no larger than 3.
3. Descend to A's other grandchildren. Terminate the search of their parent if any grandchildren is \geq A's beta. **Node B is beta-pruned**, as shown, because its value must be at least 5.
4. Once A's value is known, offer it to its parent (C) as its alpha value. So **C has alpha=3**. C will be no smaller than 3.
5. Repeat this process, descending to C's great grandchildren (D) in a depth-first fashion. **D is alpha-pruned**, because no matter what happens on its right branch, it cannot be greater than 0.
6. Repeating on E, **E is alpha-pruned** because its beta value (2) is less than its parent's alpha value (3). So no matter what happens on its right branch, E cannot have a value greater than 2.
7. Therefore **C is 3**.



Alpha-Beta Algorithm

```
function Max-Value(state, alpha, beta) returns a utility value
  if Terminal-Test (state) then return Utility(state)
  for each s in Successors(state) do
    alpha := Max (alpha, Min-Value(s, alpha, beta))
    if alpha >= beta then return beta
  end
  return alpha
```

```
function Min-Value(state, alpha, beta) returns a utility value
  if Terminal-Test (state) then return Utility(state)
  for each s in Successors(state) do
    beta := Min (beta, Max-Value(s, alpha, beta))
    if beta <= alpha then return alpha
  end
  return beta
```



Properties of Alpha-Beta Pruning

- in the ideal case, the best successor node is examined first
 - results in $O(b^{d/2})$ nodes to be searched instead of $O(b^d)$
 - alpha-beta can look ahead twice as far as minimax
 - in practice, simple ordering functions are quite useful
- assumes an idealized tree model
 - uniform branching factor, path length
 - random distribution of leaf evaluation values
- transpositions tables can be used to store permutations
 - sequences of moves that lead to the same position
- requires additional information for good players
 - game-specific background knowledge
 - empirical data

Imperfect Decisions

- complete search is impractical for most games
- alternative: search the tree only to a certain depth
 - requires a cutoff-test to determine where to stop
 - replaces the terminal test
 - the nodes at that level effectively become terminal leaf nodes
 - uses a heuristics-based evaluation function to estimate the expected utility of the game from those leaf nodes

Evaluation Function

- determines the performance of a game-playing program
- must be consistent with the utility function
 - values for terminal nodes (or at least their order) must be the same
- tradeoff between accuracy and time cost
 - without time limits, minimax could be used
- should reflect the actual chances of winning
- frequently weighted linear functions are used
 - $E = w_1f_1 + w_2f_2 + \dots + w_nf_n$
combination of features, weighted by their relevance

Example: Tic-Tac-Toe


- simple evaluation function

$$E(s) = (rx + cx + dx) - (ro + co + do)$$

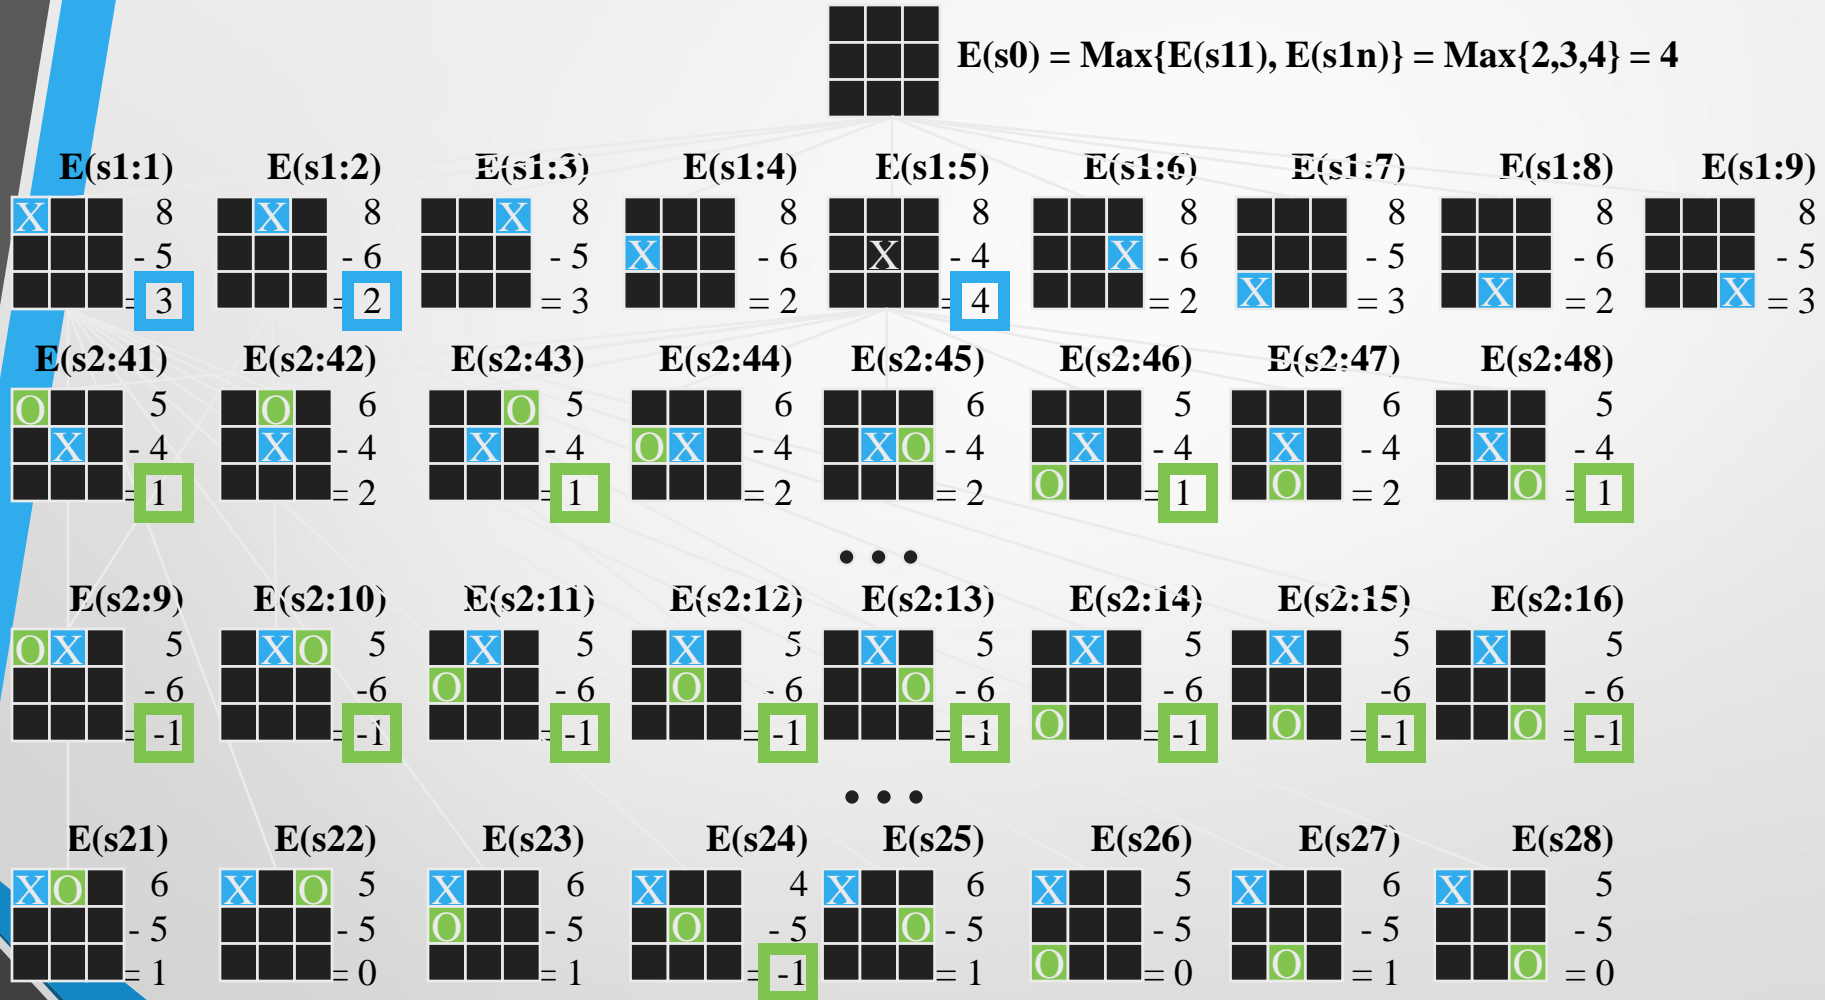
where r,c,d are the numbers of row, column and diagonal lines still available; x and o are the pieces of the two players

- 1-ply lookahead
 - start at the top of the tree
 - evaluate all 9 choices for player 1
 - pick the maximum E-value
- 2-ply lookahead
 - also looks at the opponents possible move
 - assuming that the opponents picks the minimum E-value

Tic-Tac-Toe 1-Ply



Tic-Tac-Toe 2-Ply



Checkers Case Study

- initial board configuration

- Black
 - single on 20
 - single on 21
 - king on 31

- Red
 - single on 23
 - king on 22

- evaluation function

$$E(s) = (5x_1 + x_2) - (5r_1 + r_2)$$


where

x_1 = black king advantage,



x_2 = black single advantage,

r_1 = red king advantage,

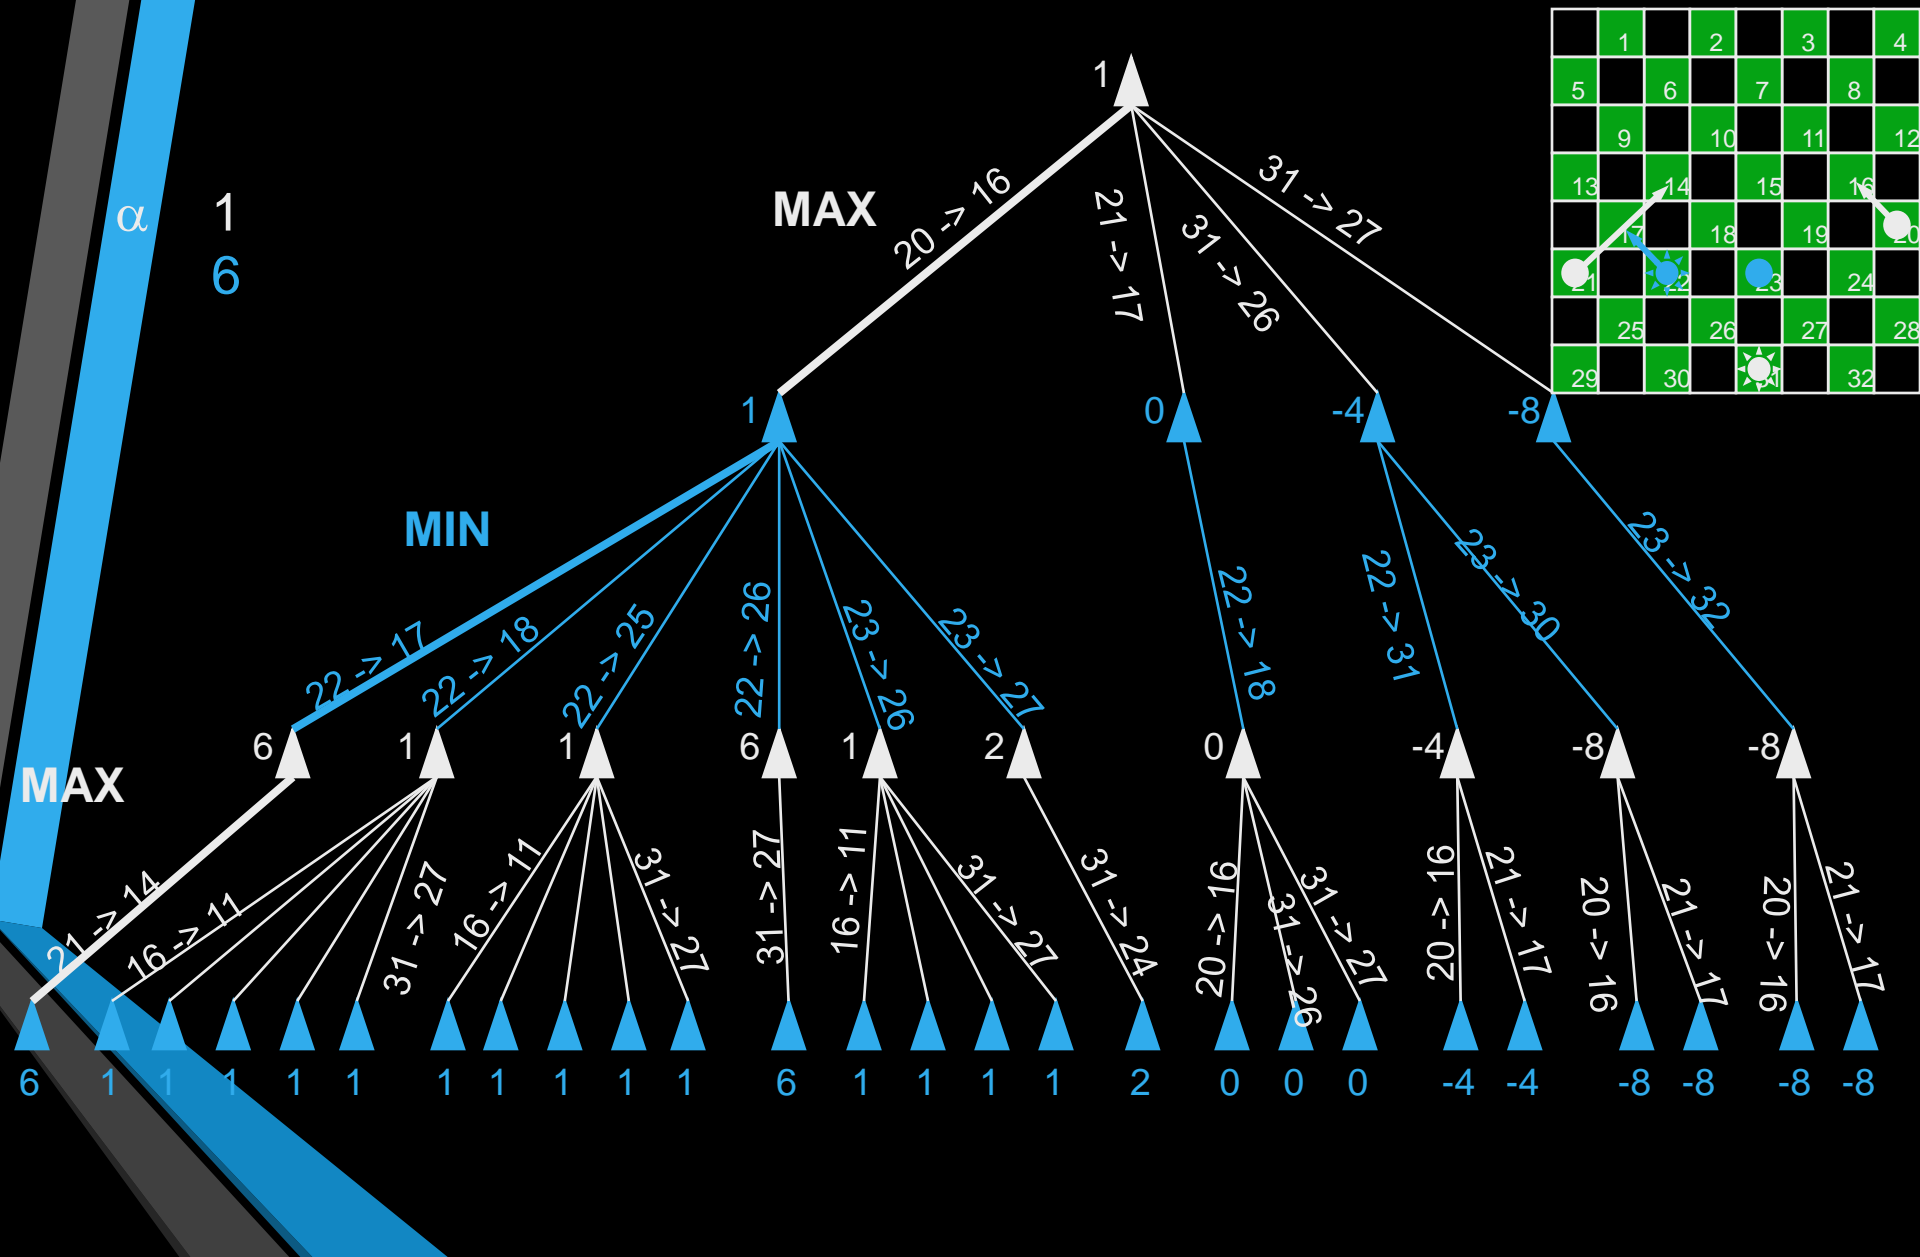
r_2 = red single advantage



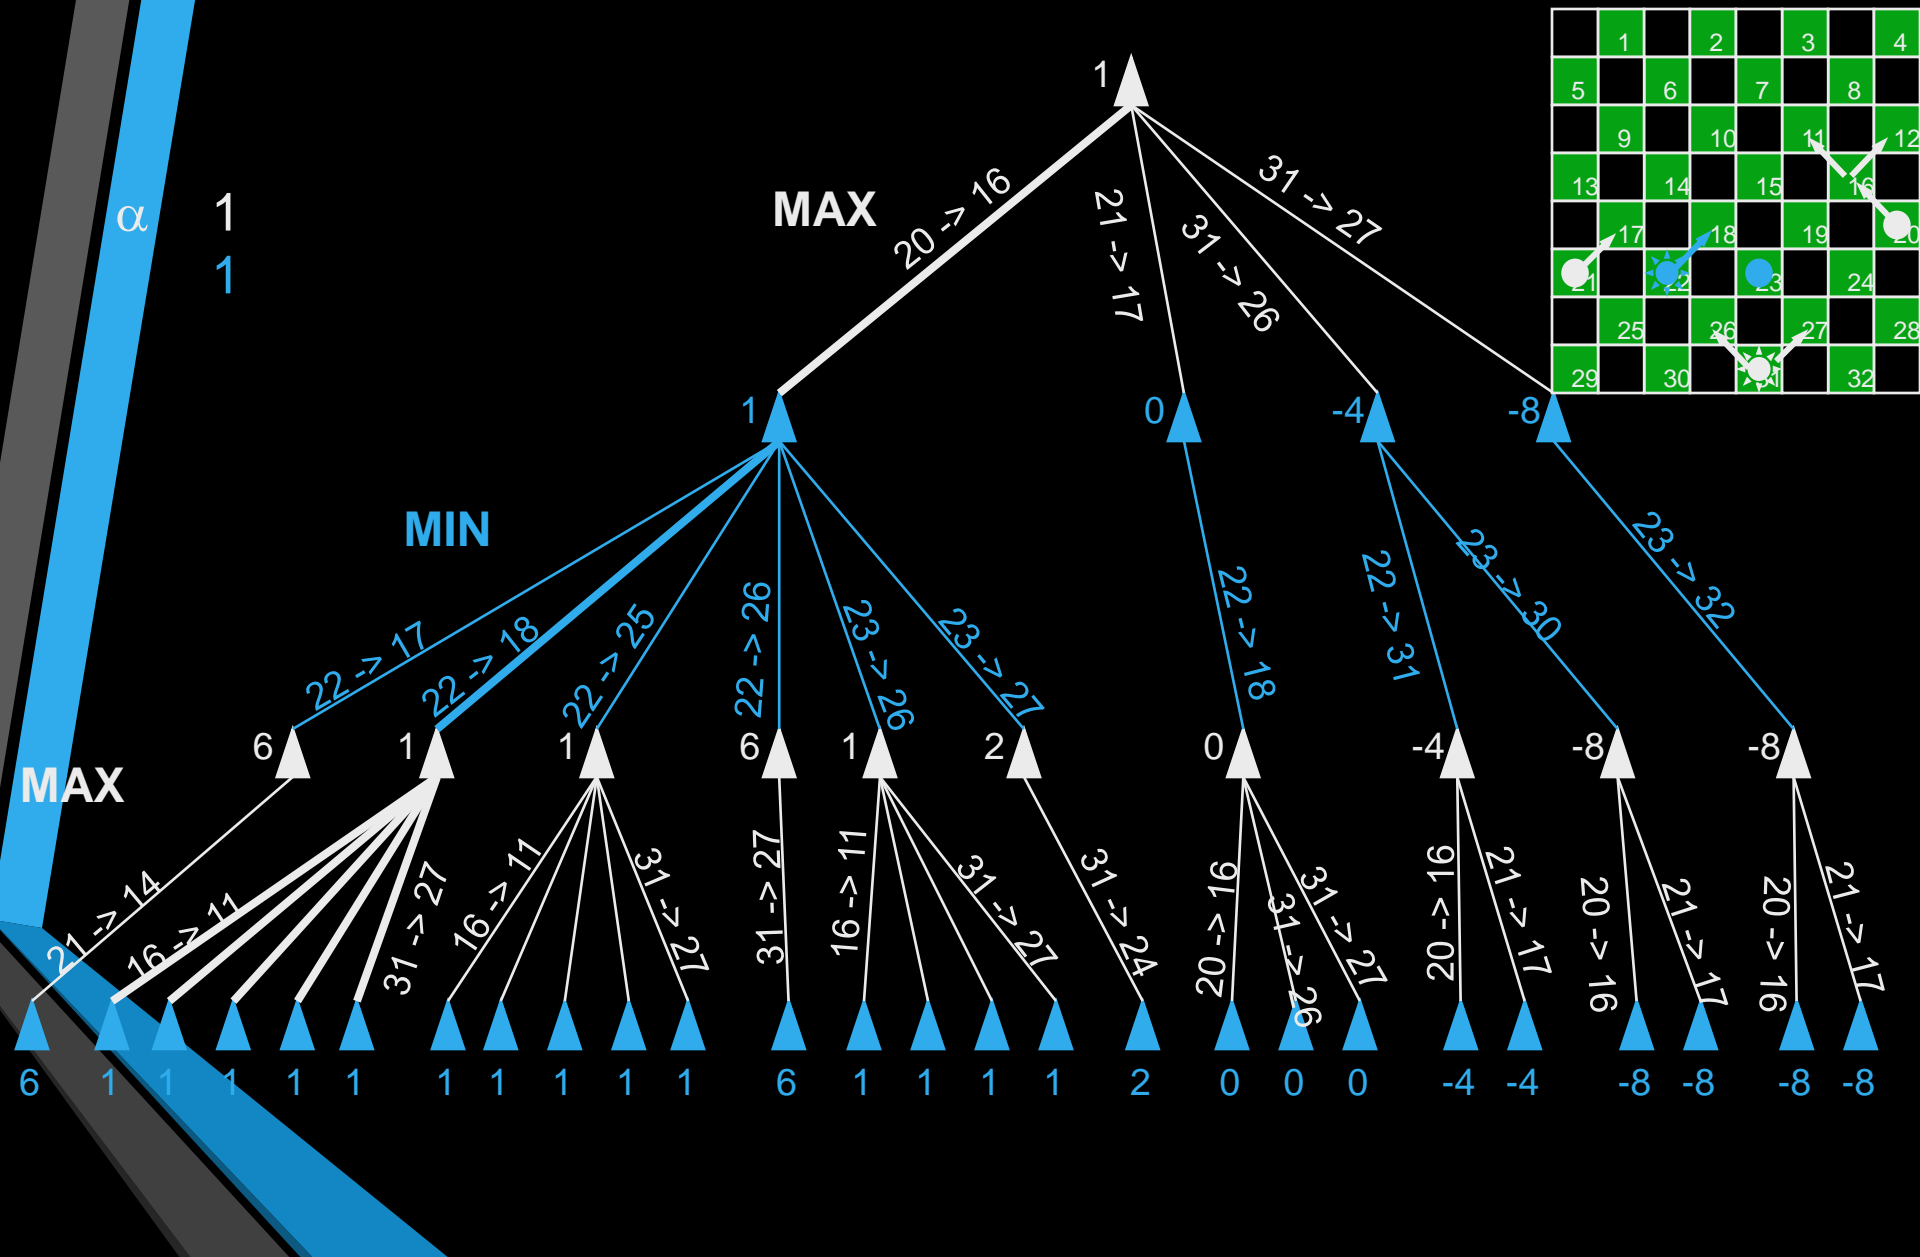
Checkers MiniMax Example



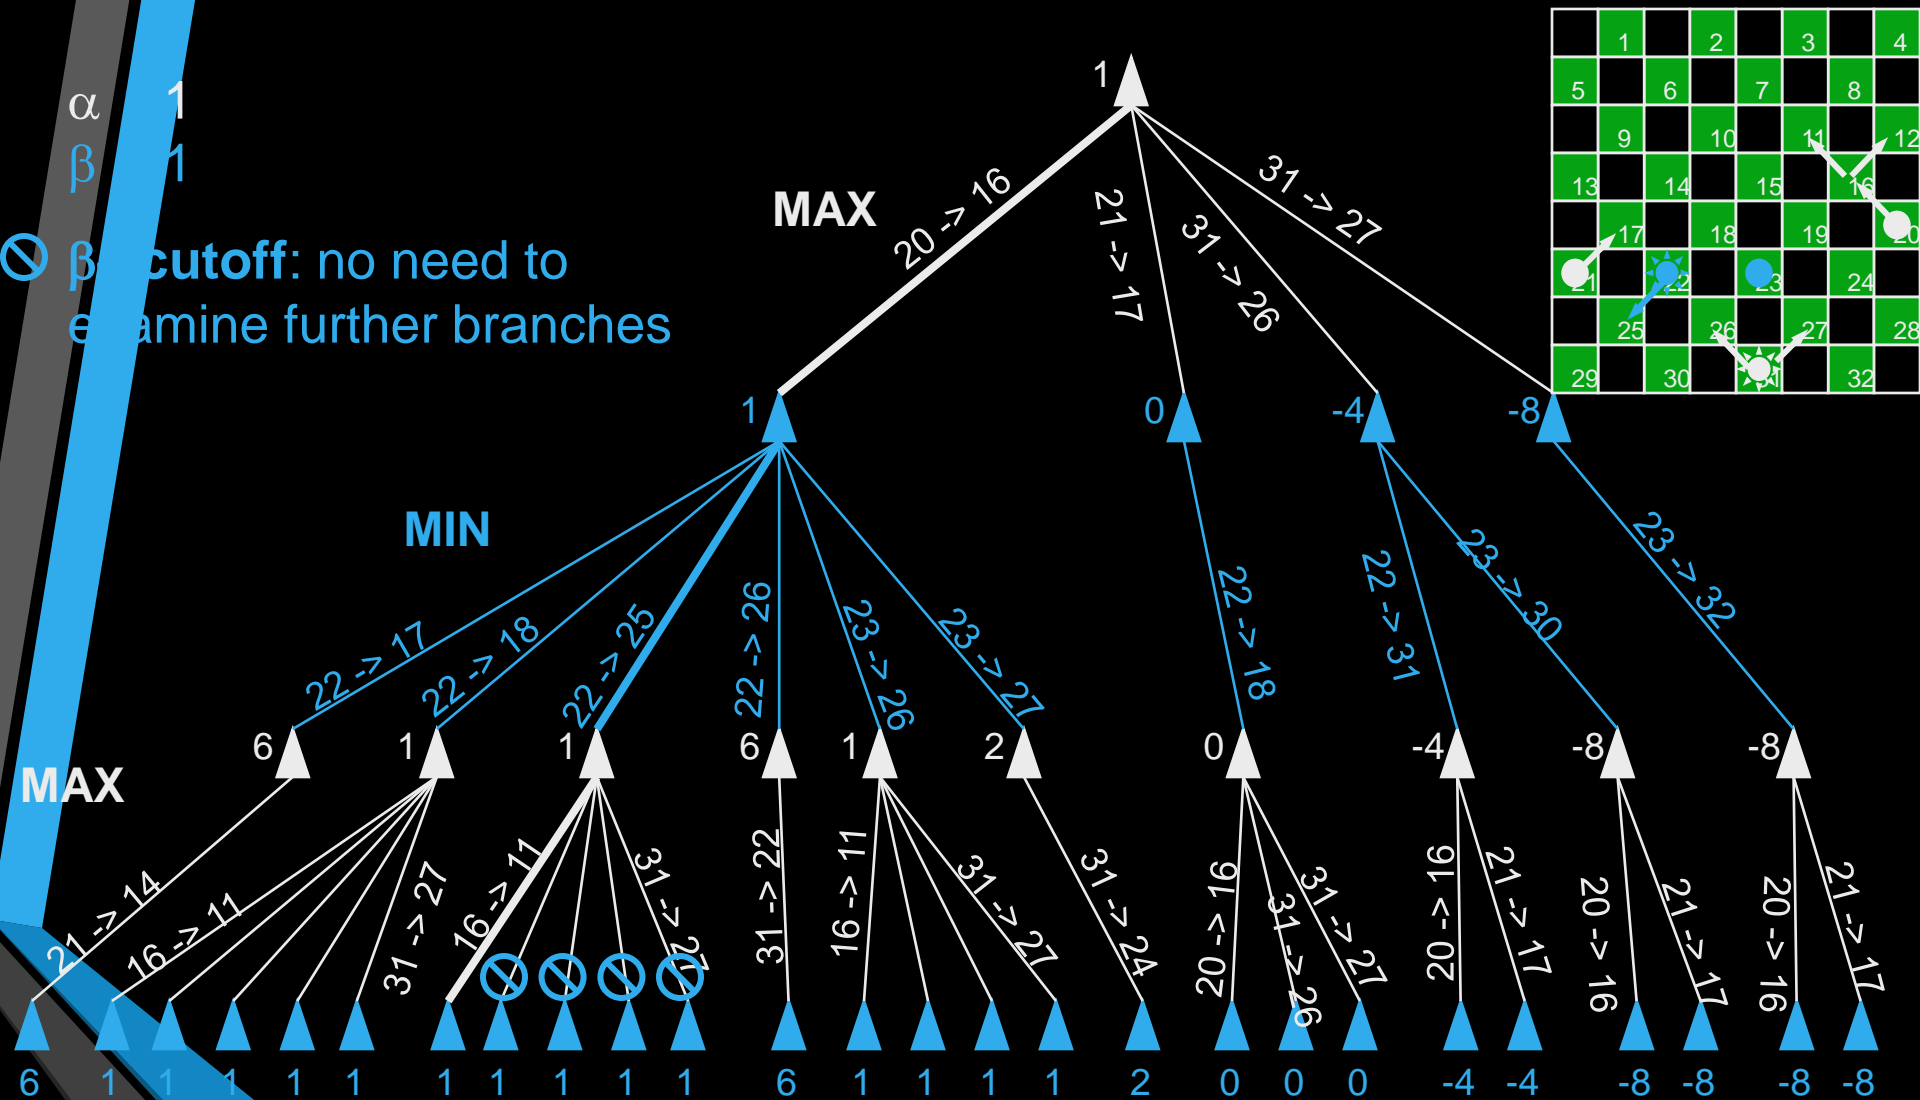
Checkers Alpha-Beta Example



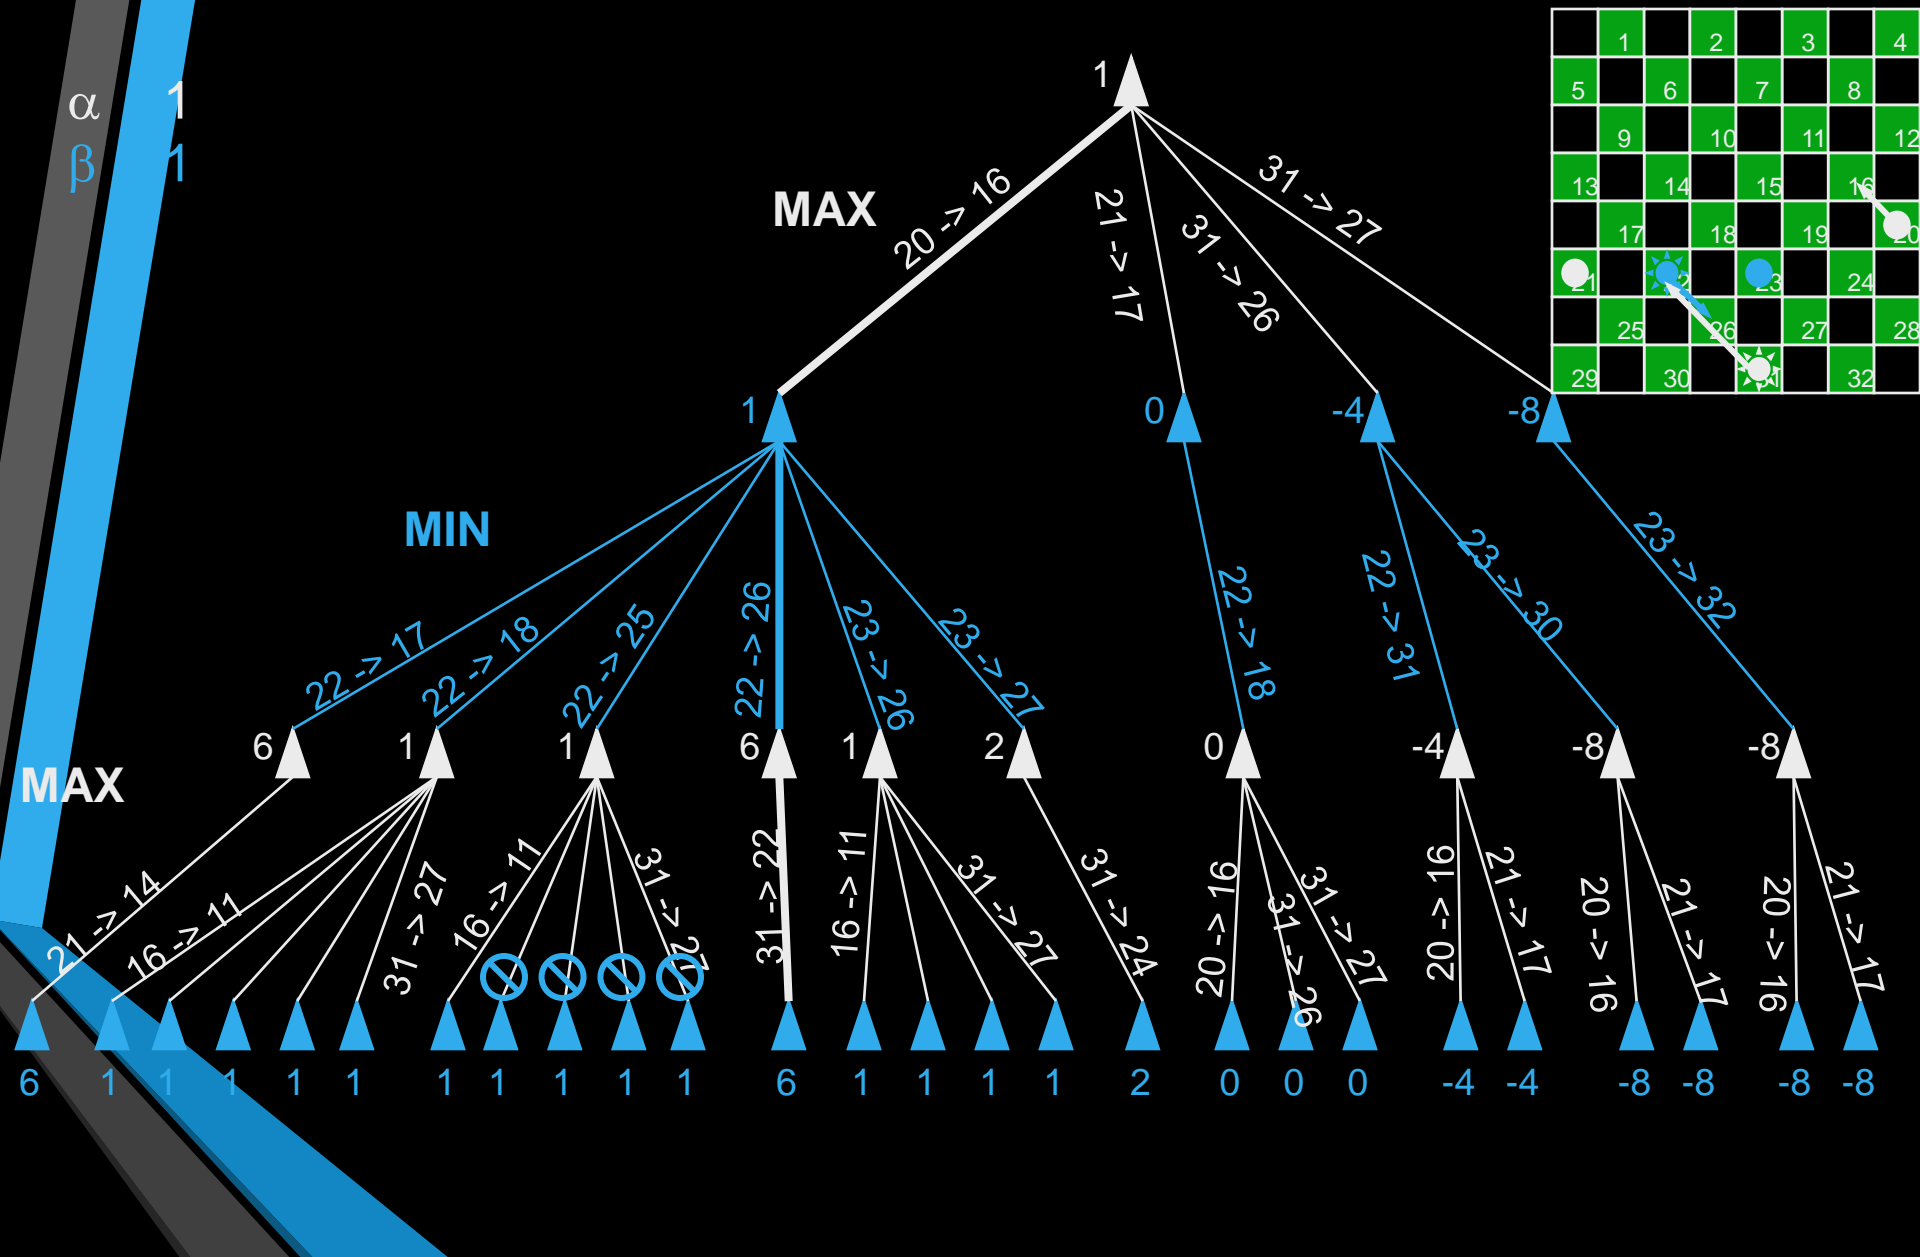
Checkers Alpha-Beta Example



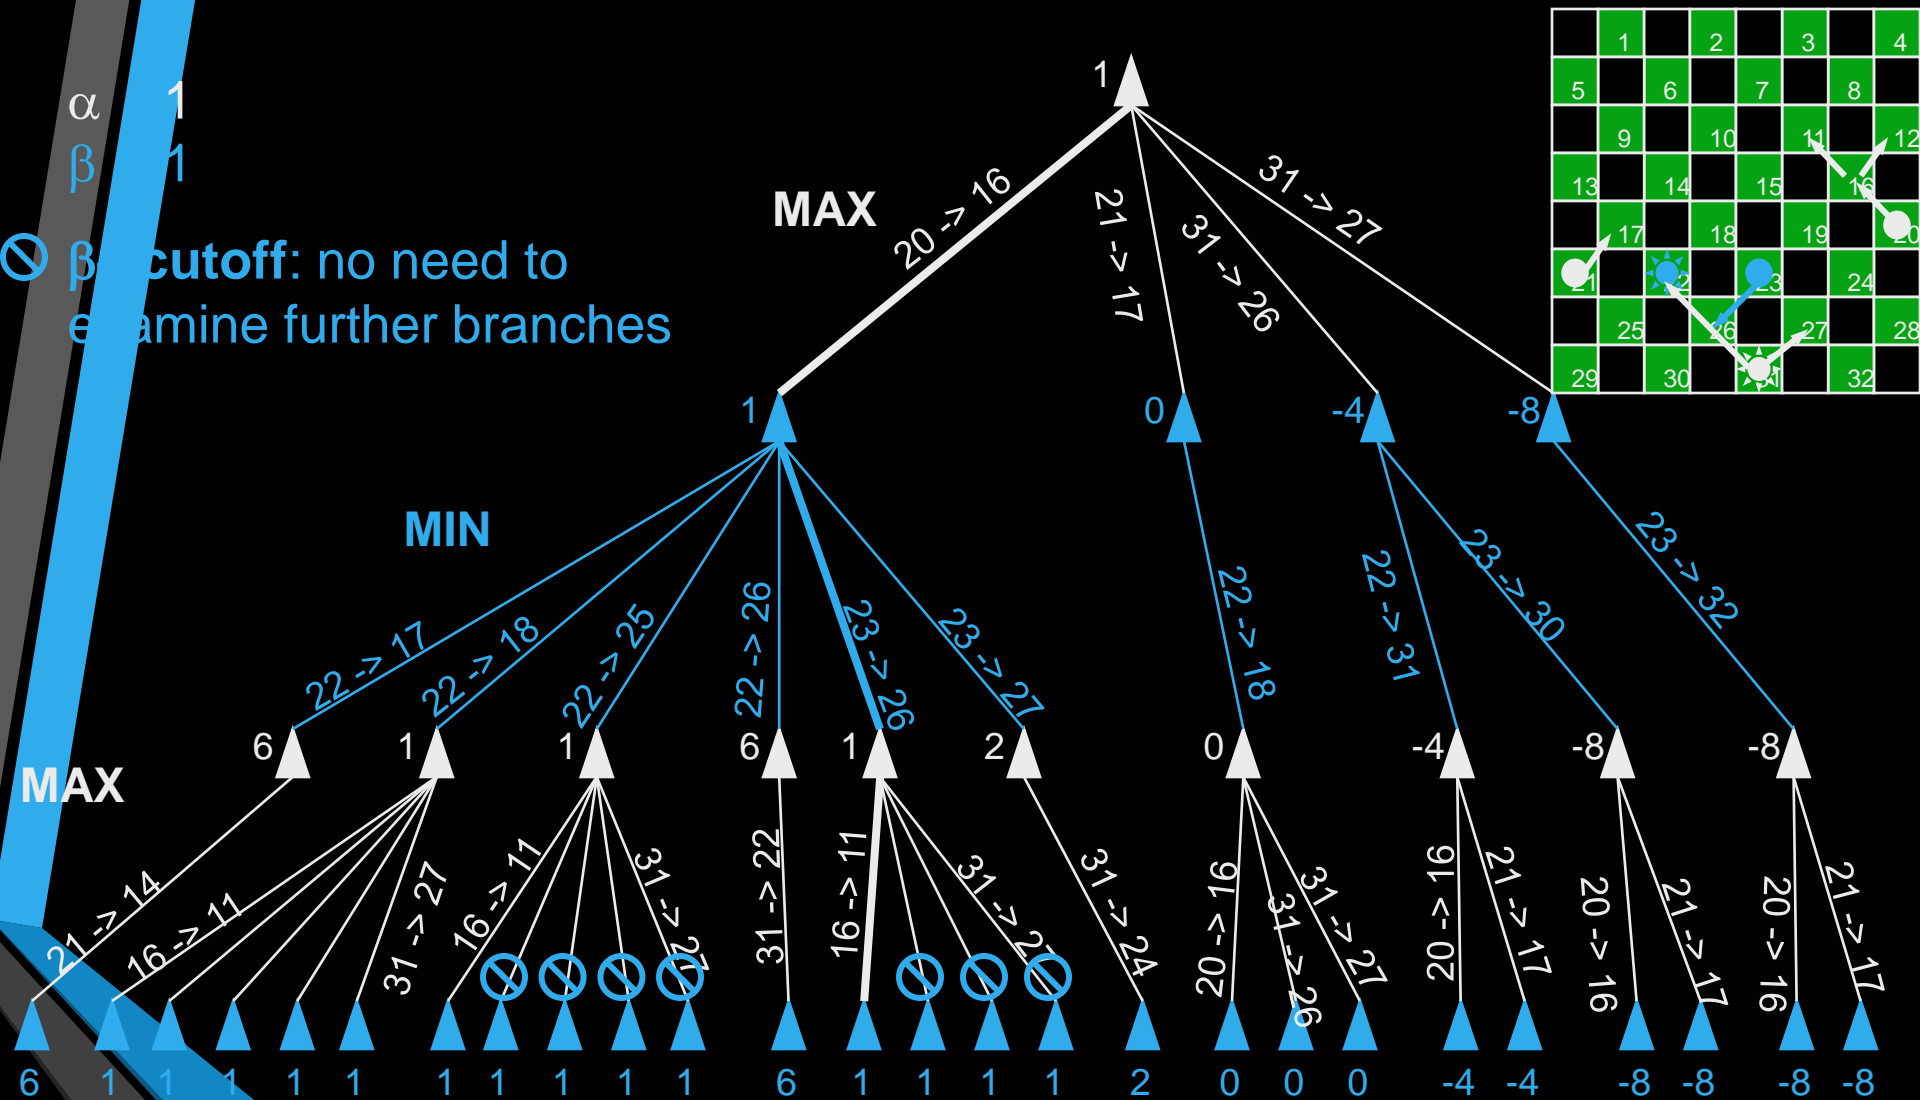
Checkers Alpha-Beta Example



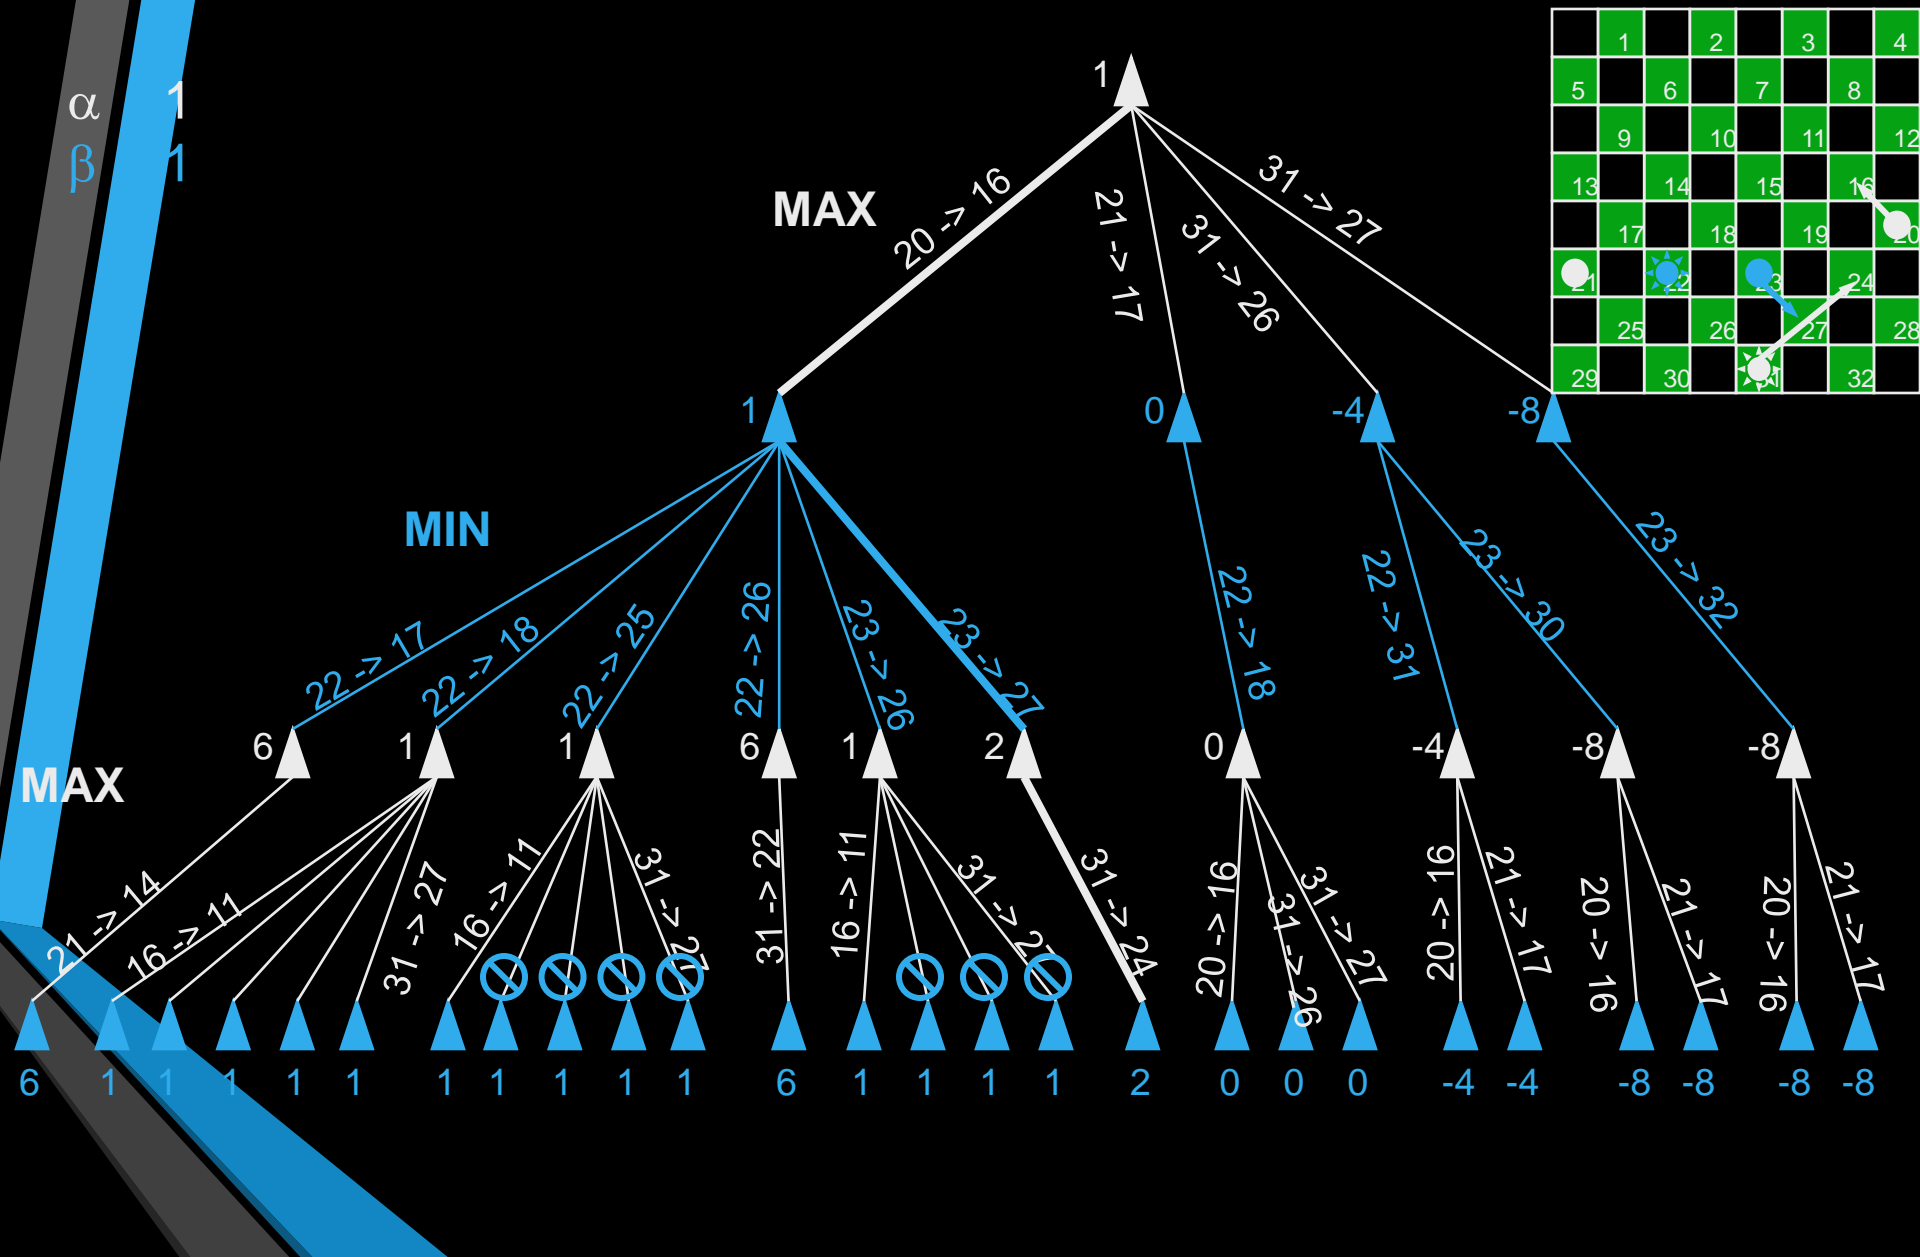
Checkers Alpha-Beta Example



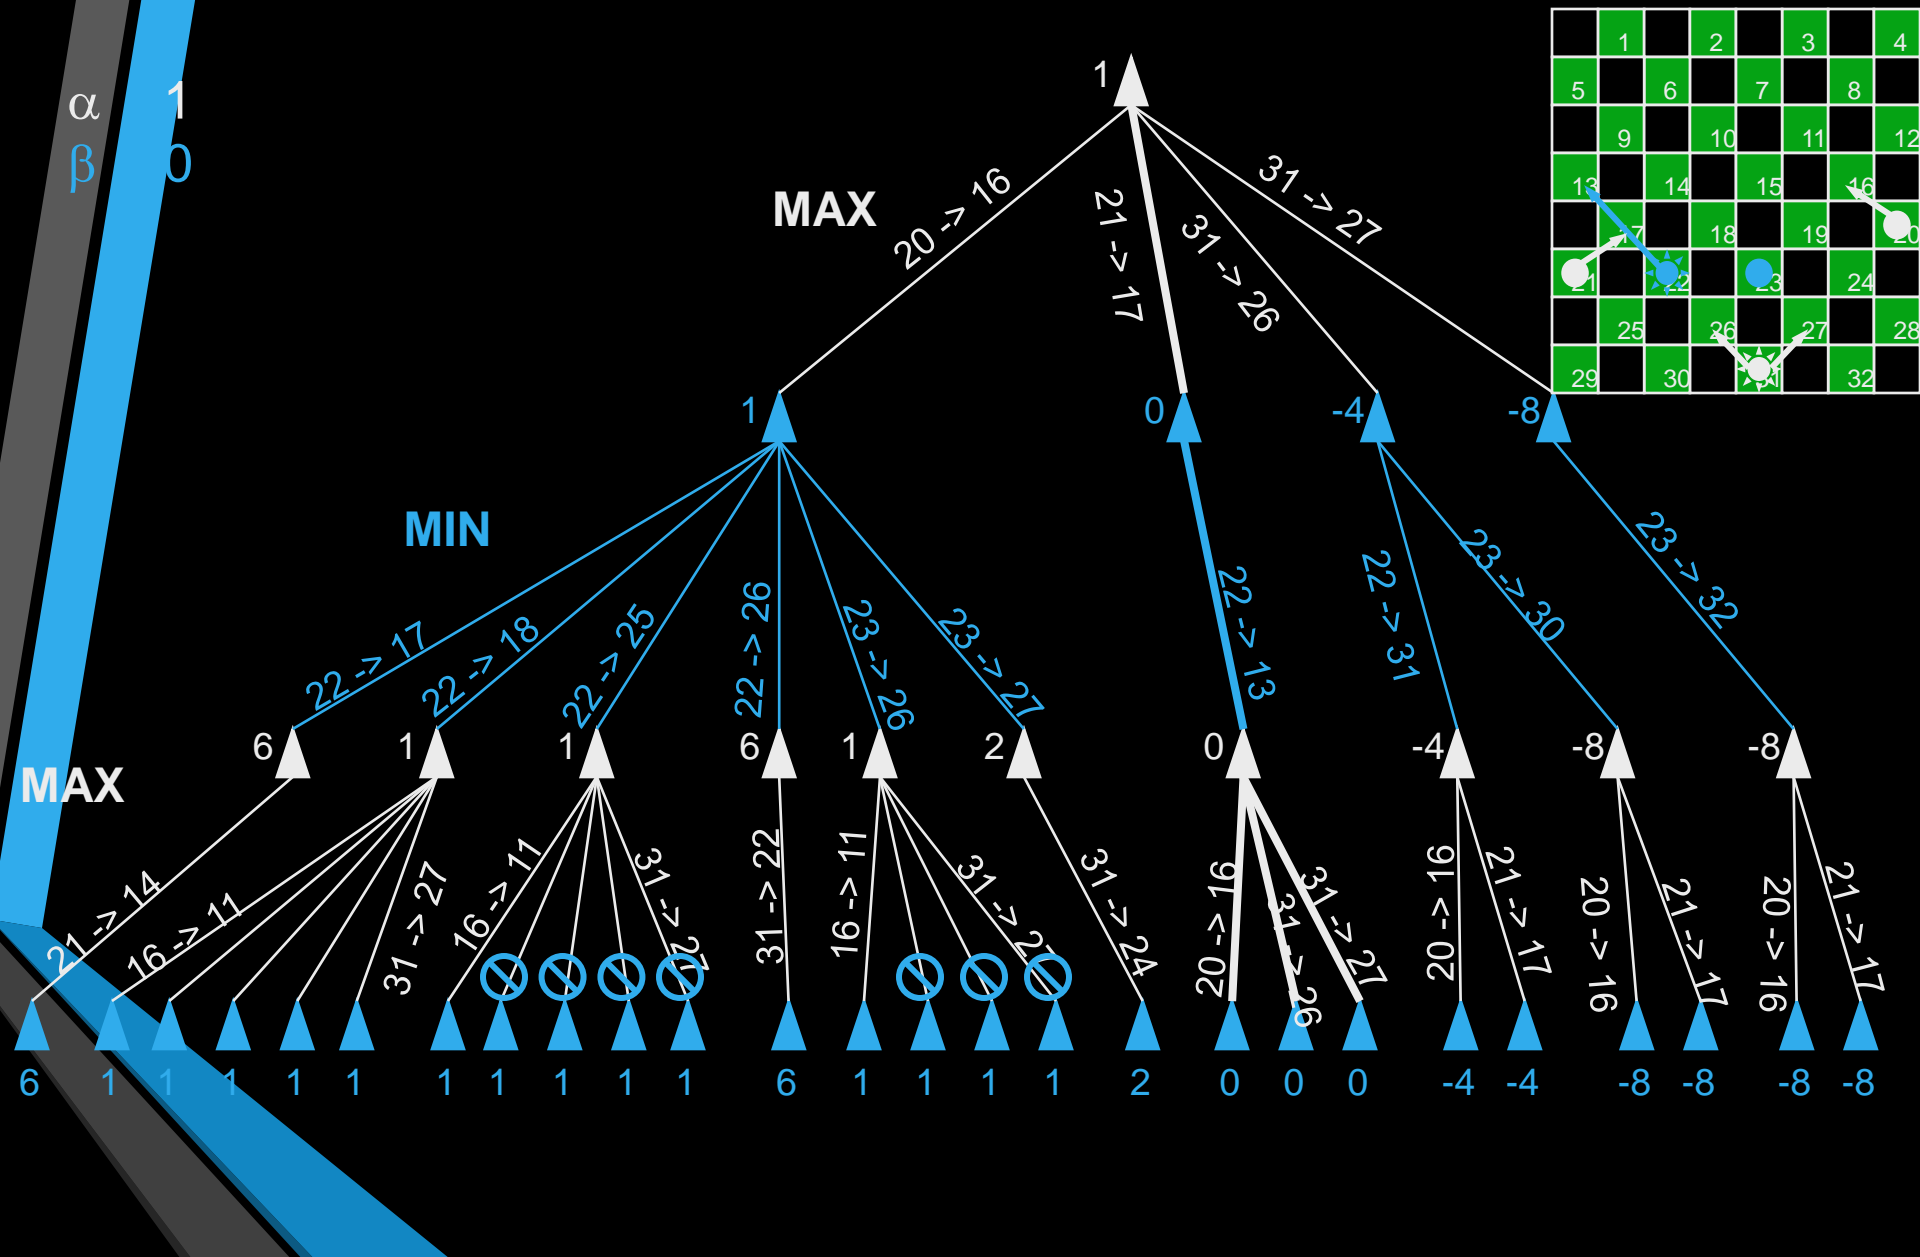
Checkers Alpha-Beta Example



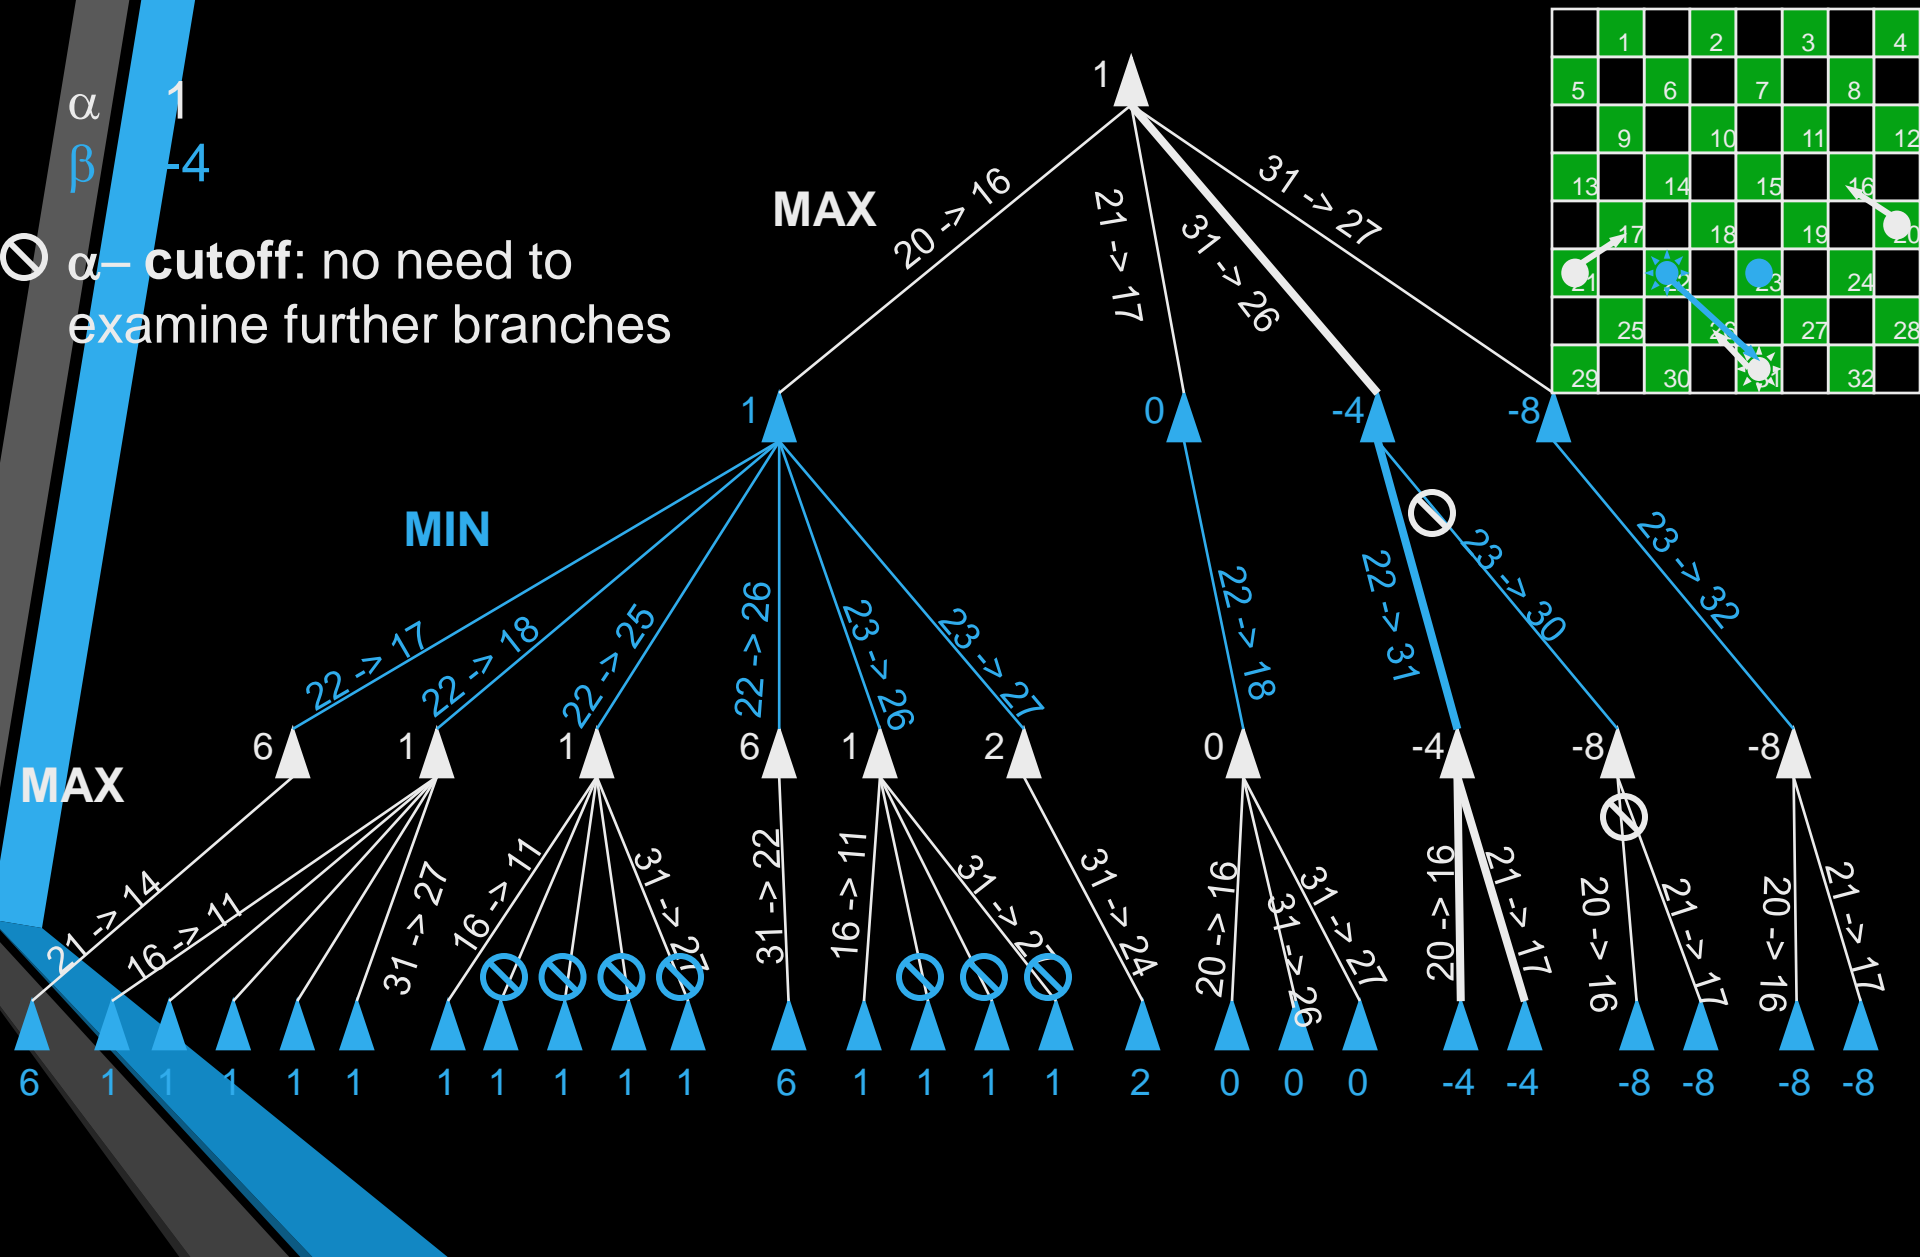
Checkers Alpha-Beta Example



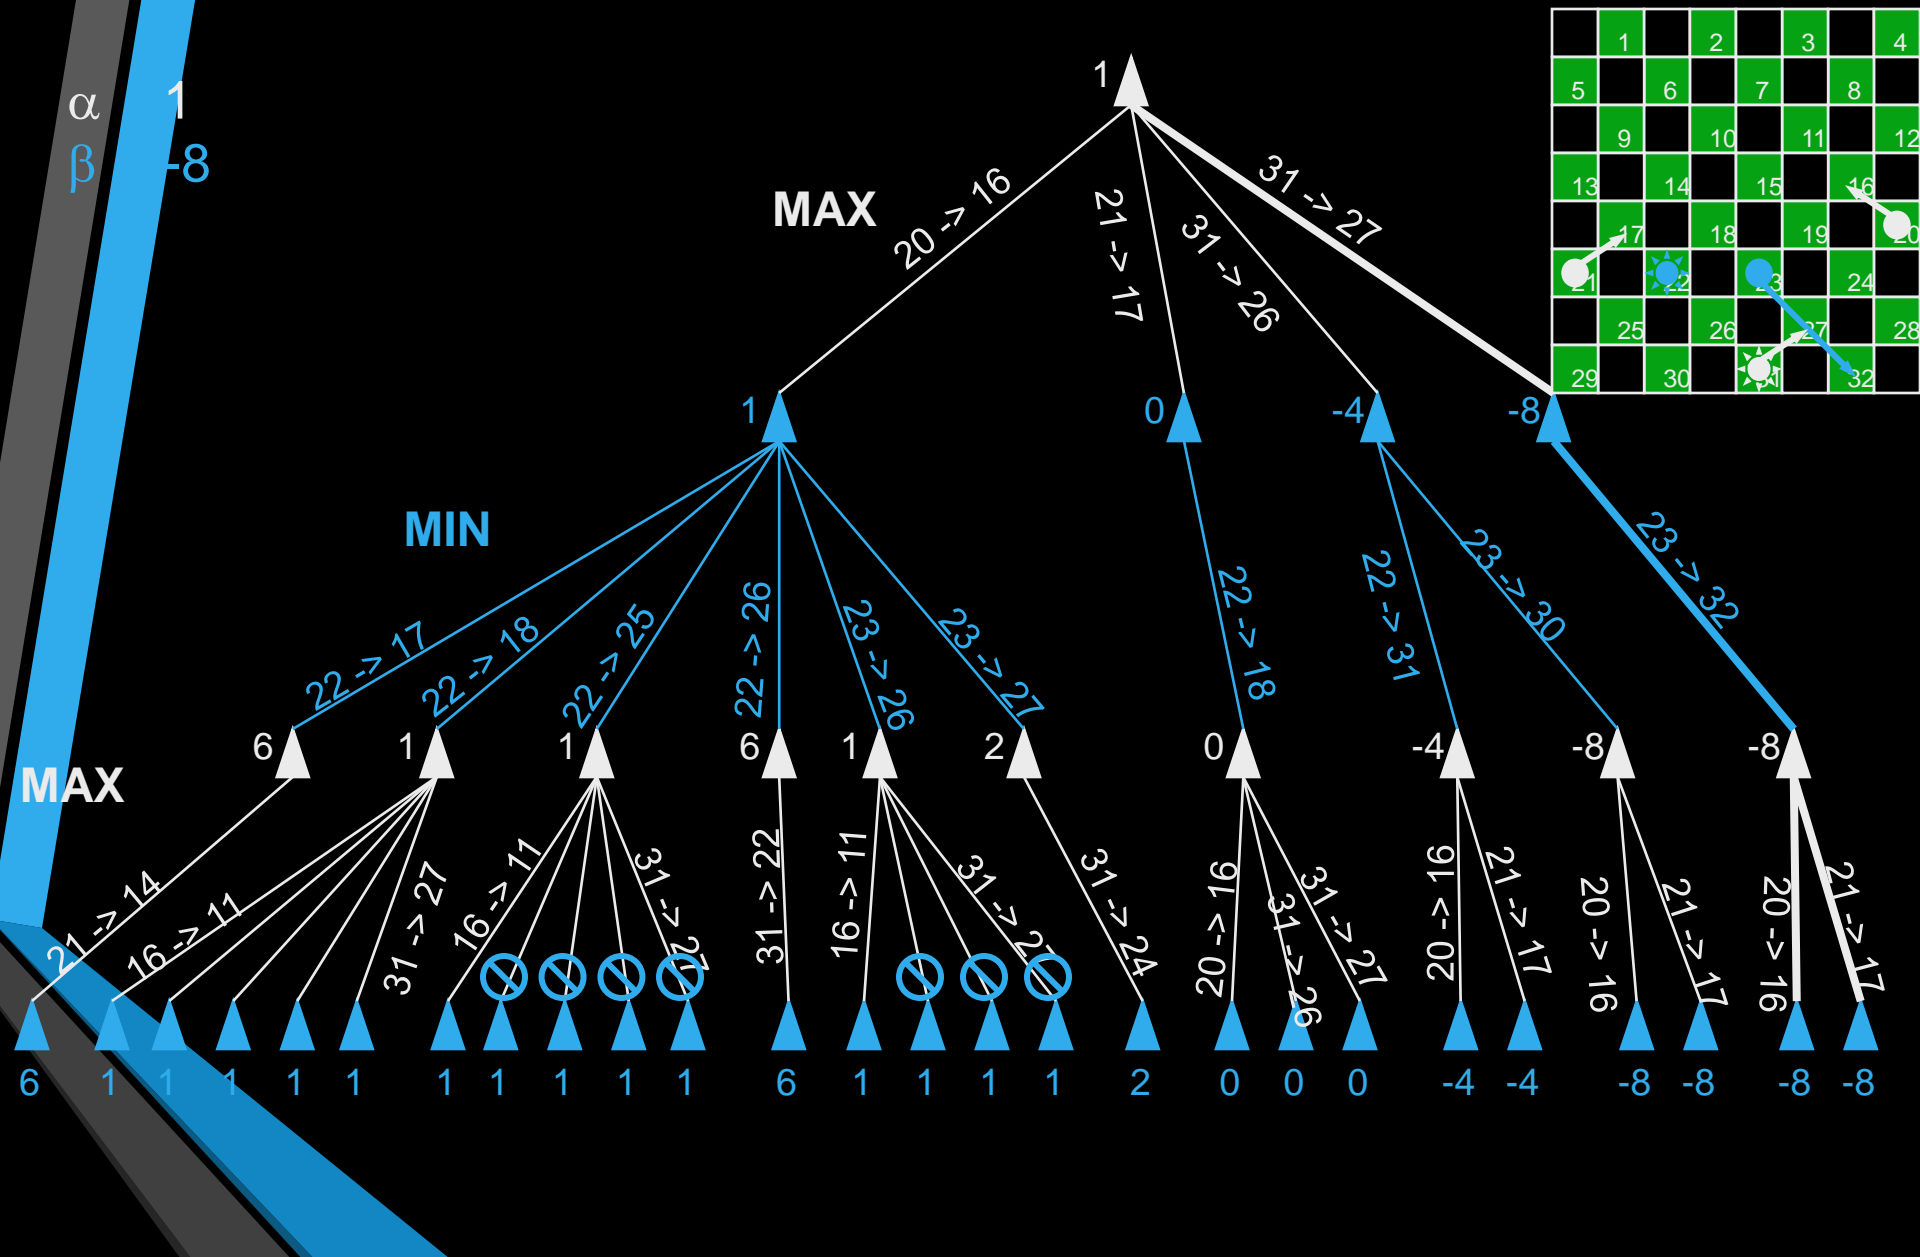
Checkers Alpha-Beta Example



Checkers Alpha-Beta Example



Checkers Alpha-Beta Example



Horizon Problem

- moves may have disastrous consequences in the future, but the consequences are not visible
 - the corresponding change in the evaluation function will only become evident at deeper levels
 - they are “beyond the horizon”
- determining the horizon is an open problem without a general solution
 - only some pragmatic approaches restricted to specific games or situation

Games and Computers

- state of the art for some game programs
 - Chess
 - Checkers
 - Othello
 - Backgammon
 - Go

Chess

- Deep Blue, a special-purpose parallel computer, defeated the world champion Gary Kasparov in 1997
 - the human player didn't show his best game
 - some claims that the circumstances were questionable
 - Deep Blue used a massive data base with games from the literature
- Fritz, a program running on an ordinary PC, challenged the world champion Vladimir Kramnik to an eight-game draw in 2002
 - top programs and top human players are roughly equal

Checkers

- Arthur Samuel develops a checkers program in the 1950s that learns its own evaluation function
 - reaches an expert level stage in the 1960s
- Chinook becomes world champion in 1994
 - human opponent, Dr. Marion Tinsley, withdraws for health reasons
 - Tinsley had been the world champion for 40 years
 - Chinook uses off-the-shelf hardware, alpha-beta search, end-games data base for six-piece positions

Othello

- Logistello defeated the human world champion in 1997
- many programs play far better than humans
 - smaller search space than chess
 - little evaluation expertise available

Backgammon

- TD-Gammon, neural-network based program, ranks among the best players in the world
 - improves its own evaluation function through learning techniques
 - search-based methods are practically hopeless
 - chance elements, branching factor

Go

- humans play far better
 - large branching factor (around 360)
 - search-based methods are hopeless
- rule-based systems play at amateur level
- the use of pattern-matching techniques can improve the capabilities of programs
 - difficult to integrate
- \$2,000,000 prize for the first program to defeat a top-level player